
VIENNA · EST. 2024 · DAILY EDITION

FinPilot

Institutional-grade analytics & decision-support infrastructure
for retail investors



The Morning Ledger — the newspaper that teaches you to read it

Retail investors are structurally left behind

87%

of retail investors underperform market benchmarks
Dalbar, 2025

**\$50–
500K**

per year — the price of institutional-grade analytics, out of reach for individuals
industry pricing

62%

of Europeans lack basic financial literacy
OECD, 2024

Why the gap persists

- Retail tools lack scientific rigor — signal bots and screeners are sold on in-sample backtests that collapse out-of-sample.
- Nobody is accountable — no provider publishes its full track record, including the misses.
- Education is disconnected from practice — textbook courses, while decisions happen in a live market.
- Advice is regulated and expensive — professional guidance stays behind private-banking walls.

THE SOLUTION

A daily research paper that grades itself — and teaches its reader



Grade Engine

Every morning at 07:45 CET the engine scans 1,800+ US stocks and issues research grades A / B / C with calibrated probability bands — e.g. “candidates with this profile moved $\geq 5\%$ within 5 days ~65% of the time.”



Open Scorecard

Every outcome is published — hits and misses alike, grade by grade. Accountability is not a promise; it is a product feature no retail competitor offers.



FinSense Academy

Each edition doubles as a lesson: concepts are taught with that morning's real, dated cases. The archive of past editions becomes the case-file library.

No buy/sell calls. No price targets. No return promises. Only grades, probabilities and plain-language reasons.

Scan → Grade → Verify → Teach — one loop, every market day

1

Scan

1,800+ US stocks screened each morning across volatility, short interest, gap, relative volume and regime factors.

2

Grade

Candidates become A / B / C research grades with a probability sentence, 2-4 factor badges and a two-sentence rationale.

3

Verify

Five days later the outcome is resolved and stamped on the public scorecard — including every miss.

4

Teach

The day's top case becomes the day's lesson; the user's own predictions feed a personal calibration score.

A

Grade A — rare, ~1 per day.
The strongest, fully-confirmed profile.

B

Grade B — strong setup.
Solid profile with one factor still unconfirmed.

C

Grade C — watch stage.
Early-tier candidates used mainly for teaching.

A production engine fed by a research lab

PRODUCTION — the scoring engine

- Volatility-aware multi-factor model: short interest, gap, relative volume, ATR and market regime — each factor validated point-in-time.
- Point-in-time data contracts — no look-ahead: every input is stamped with what was knowable that morning.
- Execution-replay infrastructure: triple-barrier P&L with realistic costs, slippage and commissions.
- Calibrated probability bands per grade, recomputed as the archive grows.

RESEARCH LABS — what feeds it next

- Multi-agent deep reinforcement learning line (PPO, SAC, TD3, recurrent agents) for adaptive, regime-specialised decision models.
- Shadow-validation pipeline: candidate models run silently beside production until they beat it on locked data.
- Agent workflows: autonomous research agents that generate, test and document factor hypotheses.
- Privacy-first local LLM inference for personal analysis — user data never leaves the device.

Research graduates to production only after it clears the validation bar — never on a promising backtest alone.

The discipline is the technology

	Typical retail tools	FinPilot standard
Validation	In-sample backtests, tuned until they look good	Locked out-of-sample tests; walk-forward audits; cluster-bootstrap statistics
Performance claims	“Accuracy” on favorable price moves	Triple-barrier execution replay with realistic costs — precision ≠ realized P&L, and we measure both
Track record	Wins showcased, losses deleted	Public scorecard of every graded call, misses included
Signals	Static indicators (RSI, MACD) recycled since the 1980s	Regime-aware factor research; our own tests retired RSI/MACD — and our own legacy score — when they failed
Learning	Static courses, disconnected from the market	Lessons auto-generated from real, dated, pre-outcome-graded cases

We attack our own results before the market can

30+

One research cycle: 30+ independent tests across factors, exits, costs and universes before any production change.

BACKTEST BATTERY

5,000+

An archive of graded, dated, pre-outcome calls — the raw material for calibration and for FinSense case files.

RESOLVED SIGNALS

6 / 6

The volatility edge held in every month tested, after realistic costs — our most robust, repeatedly confirmed finding.

MONTHS STABLE

0

Findings that failed re-testing were retired — including our own legacy scoring formula. The jury sees what survived.

VANITY METRICS

Honesty as strategy: the platform publishes the same rigorous, cost-adjusted numbers to users that it shows to investors.

Compliance is architecture, not an afterthought



No advice language

The user interface contains no buy/sell recommendations, no target prices, no position sizing and no return promises — by design, on every surface. FinPilot is decision-support software, not licensed investment advice.



Full auditability

Every published grade carries a timestamp and an engine-configuration fingerprint; the archive is immutable. Any output can be reproduced and audited after the fact.



Privacy-first AI

GDPR-compliant foundations today; roadmap to local LLM inference so personal financial analysis runs on the user's device — data never leaves it.



Regulatory roadmap

Built to remain outside licensed-advice activity today, with a prepared path toward MiFID-compliant personalised guidance as the product deepens — compliance grows with the vision.

A financial literacy lab that improves itself

Like a language-learning app for markets: daily micro-lessons, streaks and quests — but every exercise is built from this morning's real market and, later, from the stocks the user actually follows.

Daily quests

5-minute tasks tied to the morning edition — read a case, predict an outcome, earn your streak. Learning happens inside the live market, not beside it.

Calibration trainer

Users grade candidates before seeing the engine's grade; a personal calibration score shows exactly where their judgement is over- or under-confident.

Blind-spot map

The platform learns from each user's decisions and mistakes, maps their weak areas, and steers the next lessons there — personal financial analysis in service of learning.

Self-improving content

Lessons are auto-generated from the growing archive of resolved cases: every example is real, dated and graded before the outcome was known. More market days → better lab.

12 domains · 2 learning tracks · 120+ glossary terms · EN / DE / TR — closing Europe's financial literacy gap, one edition a day.

A \$12B market — entered where trust matters most

\$12B

retail trading & investing tools market, growing
~18% CAGR

150M+

self-directed retail investors in Europe alone

200M+

Europeans who want to learn — FinSense free tier
is the wedge

1

Europe first

Launch from Vienna into DACH & EU: highest literacy need, EU literacy agenda tailwind, GDPR-native design as a trust advantage. US market data already powers the engine.

2

United States

The largest self-directed investor base; product is US-market-native from day one, so entry is distribution, not rebuild.

3

Global, multi-market

Multi-market engine expansion (EU exchanges, then global), 5+ languages, B2B white-label as the scaling channel.

Freemium — trust first, revenue follows

Free — Daily Brief

€0

- Morning brief with 1-2 graded candidates (web + Telegram)
- Open scorecard — full transparency for everyone
- FinSense core lessons & daily quests

Premium — Full Edition

€19–29 / mo

- Complete daily edition: all graded candidates + risk notes
- Personal calibration trainer & blind-spot map
- Full FinSense tracks, archive & case-file library

B2B — API & White-label

custom

- Graded-research feed for neobrokers & fintechs
- Co-branded FinSense education modules
- The future channel for multi-market scale

80%+ gross margin (SaaS) · education-led funnel lowers acquisition cost and churn

From morning paper to personal financial advisor & educator

TODAY — MVP

- The Morning Ledger: daily graded US-market editions
- Open scorecard + validated factor engine
- FinSense lessons from real archived cases
- Working prototype; multi-agent DRL in the lab

THE PRODUCT WE ARE BUILDING TOWARD

- A professional-grade investing experience for individuals — the quality of service today reserved for private-banking clients
- Personal AI that analyses each user's finances, decisions and learning — and coaches them to their best outcomes
- Multi-agent AI working for the user: research, monitoring and workflow agents around a compliance core
- Multi-market coverage: Europe → US → global

Every euro of this project funds the bridge between the two columns.

THE ASK

aws Preseed – Deep Tech · €267,000

WP1	Truth Engine & data integrity — live public scorecard	€35K
WP2	Scoring R&D: shadow validation + multi-agent DRL lab	€85K
WP3	Product build-out: Morning Ledger + Reader's Desk	€70K
WP4	FinSense impact engine: quests, calibration, auto-lessons	€55K
WP5	Regtech, compliance & security layer	€35K
WP6	Market validation: beta launch, 1,000 users, payments	€50K

18-month project · total ~€330K · ≥10% shareholder equity secured by founder

MILESTONES

- M4 — public scorecard live
- M10 — Morning Ledger public launch
- M12 — shadow-validated V2 engine in production
- M18 — 1,000 beta users; Seedfinancing – Deep Tech follow-up